

Administrivia

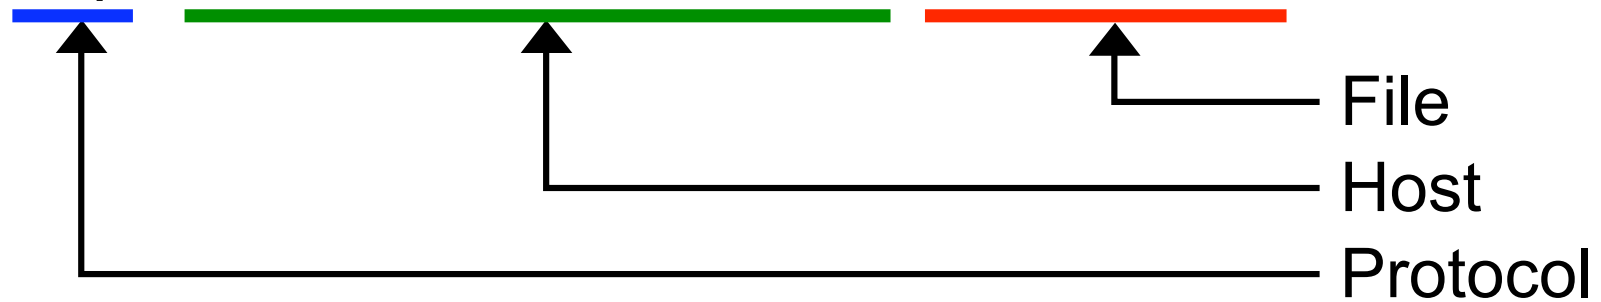
- **Lab 2 due right now**
 - Free extension to midnight for being here
 - Put `/* Attended-Lecture */` at top of `reliable.c`
- **Midterm exam one week from today**
 - Open Book, Open notes, no electronic devices allowed
 - Feel free to print out and bring lecture slides
- **SCPD students:**
 - Email `cs144-staff@scs.stanford.edu` with your exam monitor information
 - Please ensure the email subject is “exam monitor”
- **Any other students with special exam needs**
 - Please email `cs144-staff` to make arrangements

Outline

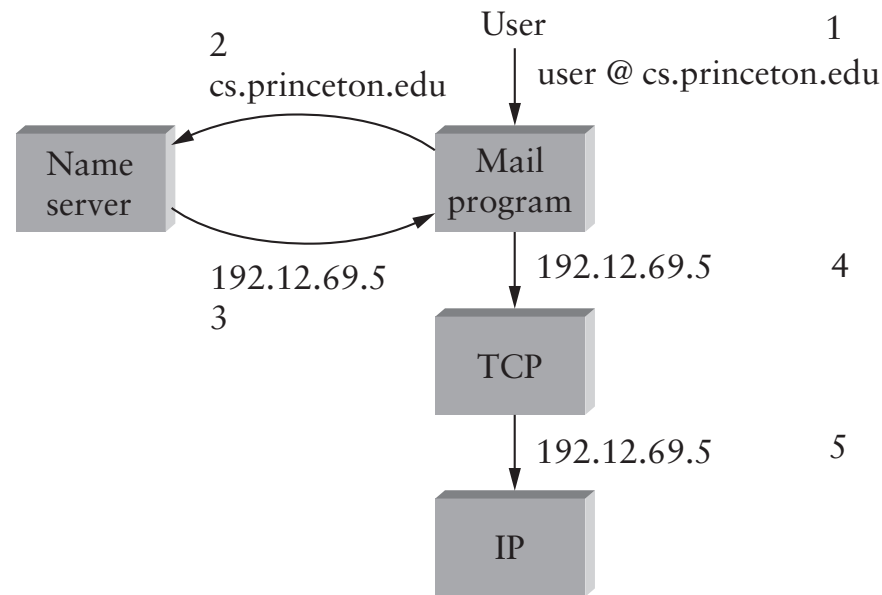
- DNS architecture
- DNS protocol and resource records (RRs)
- Record types: A, NS, glue, MX, SOA, CNAME
- Reverse lookup
- Load balancing
- DNS security

Parsing a URL

http://cs144.scs.stanford.edu/labs/sc.html



Motivation



- **Users can't remember IP addresses**
 - Need to map symbolic names (`www.stanford.edu`) → IP addr
- **Implemented by library functions & servers**
 - `getaddrinfo ()` talks to server over UDP
- **Actually, more generally, need to map symbolic names to values**

hosts.txt system

- **Originally, hosts were listed in a file, hosts.txt**
 - Email global network administrator when you add a host
 - Administrator mails out new hosts.txt file every few days
- **Would be completely impractical today**
 - hosts.txt today would be huge (gigabytes)
 - What if two people wanted to add same name?
 - Who is authorized to change address of a name?
 - People need to change name mappings more often than every few days (e.g., Dynamic IP addresses)

Goals of DNS

- **Scalability**

- Must handle huge number of records
- Potentially *exponential* in name size—because custom software may synthesize names on-the-fly

- **Distributed control**

- Let people control their own names

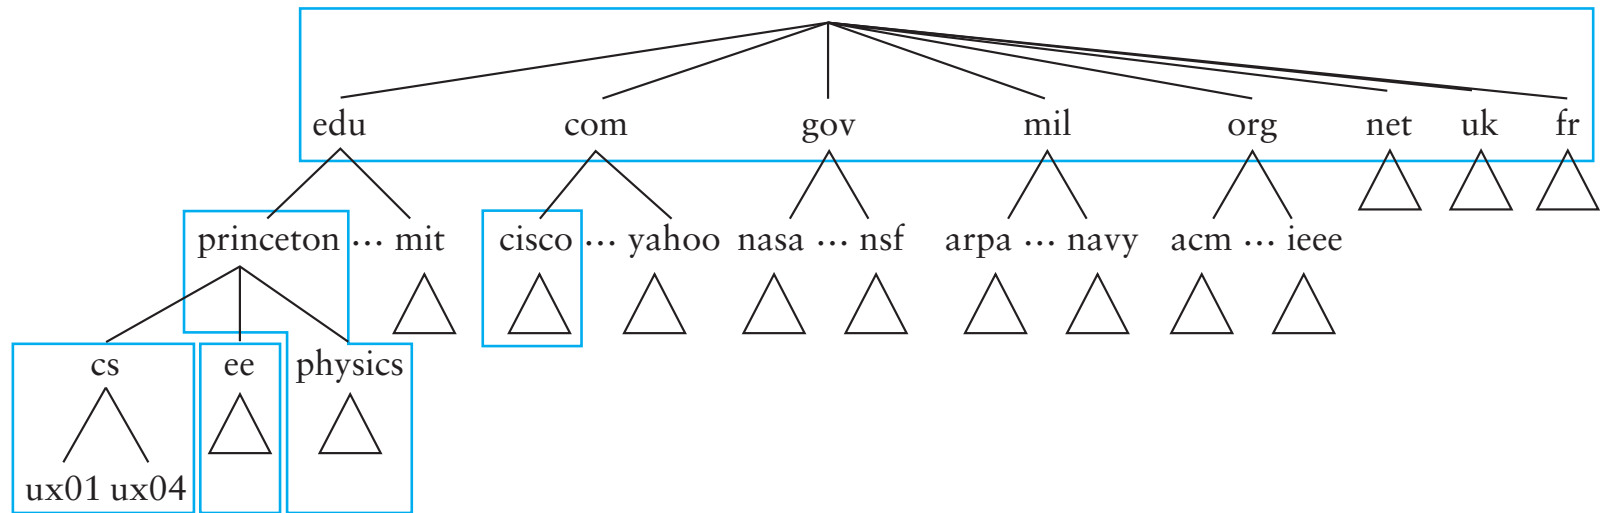
- **Fault-tolerance**

- Old software assumed `hosts.txt` always there
- Bad potential failure modes when name lookups fail
- Minimize lookup failures in the face of other network problems

The good news

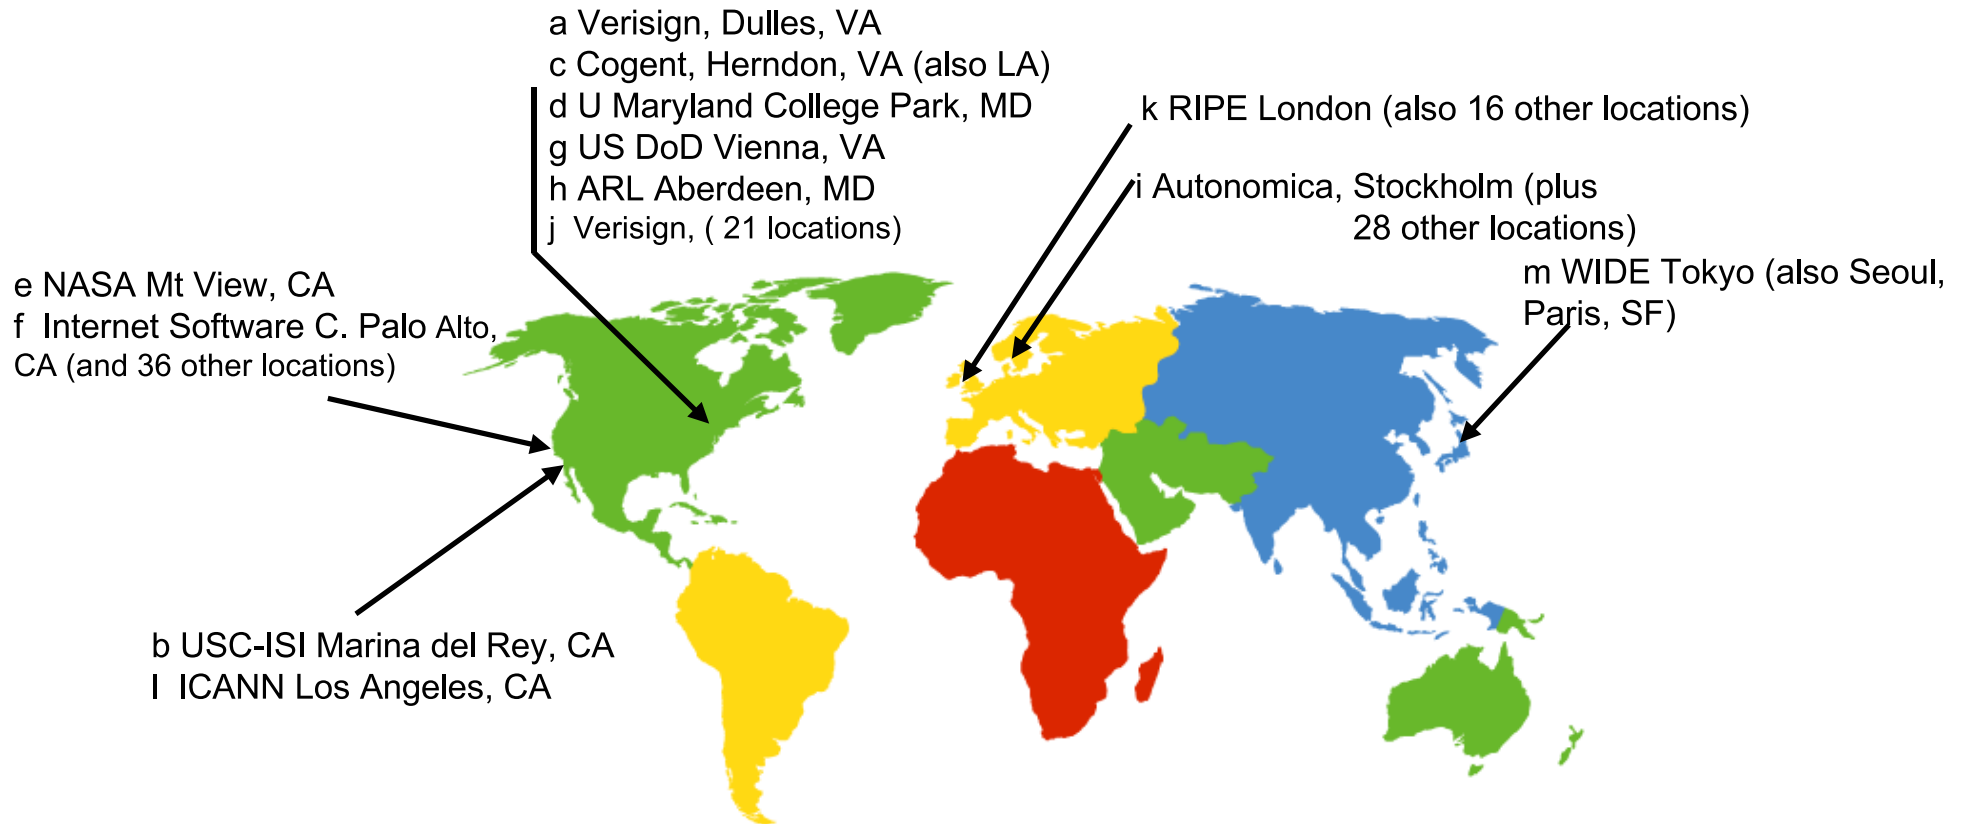
- **Properties that make DNS goals easier to achieve:**
 1. **Read-only or read-mostly database**
 - People typically look up hostnames much more often than they are updated
 2. **Loose consistency**
 - When adding a machine, may be okay if info takes minutes or hours to propagate
- **These suggest approach w. aggressive caching**
 - Once you have looked up hostname, remember result
 - Don't need to look it up again in near future

Domain Name System (DNS)



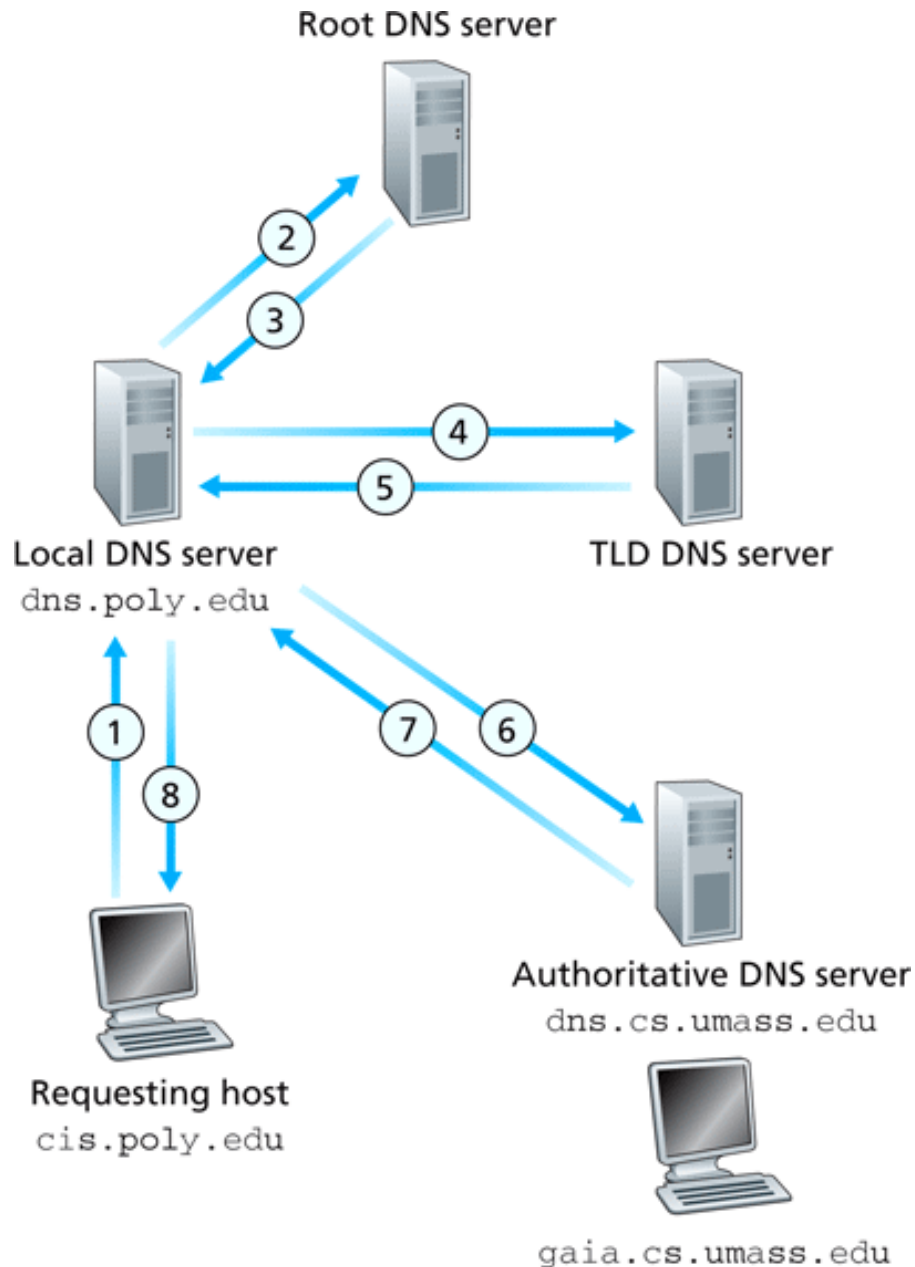
- **Break namespace into a bunch of zones**
 - . (“root”), edu., stanford.edu., cs.stanford.edu., ...
 - Zones separately administered \Rightarrow **delegation**
 - Parent zones tell you how to find servers for subdomains.
- **Each zone served from several replicated servers**

Root servers



- **Root (and TLD) servers must be widely replicated**
 - For some, use various tricks like IP anycast

DNS software architecture



- Two types of query
 - Recursive
 - Non-Recursive
- Apps make **recursive** queries to local DNS server (1)
- Local server queries remote servers **non-recursively** (2, 4, 6)
 - Aggressively caches result
 - E.g., only contact root on first query ending `.umass.edu`

DNS protocol

- **TCP/UDP port 53**
- **Most traffic uses UDP**
 - Lightweight protocol has 512 byte UDP message limit
 - retry w. TCP if UDP fails (e.g., reply truncated)
- **TCP requires message boundaries**
 - Prefix all messages w. 16-bit length
- **Bit in query determines if query is recursive**

Resource records

- All DNS info represented as resource records (RR):

name [TTL] [class] type rdata

- *name* – domain name (e.g., `www.stanford.edu.`)
- *TTL* – time to live in seconds
- *class* – for extensibility, usually IN (1) “Internet”
- *type* – type of the record
- *rdata* – resource data dependent on the *type*

- Two important DNS RR types:

- *A* – Internet address (IPv4)
- *NS* – name server

- Example resource records (`dig stanford.edu`):

```
stanford.edu.      3600   IN  A    171.67.216.4
stanford.edu.      3600   IN  A    171.67.216.7
stanford.edu.      6171   IN  NS   Argus.stanford.edu.
...
```

Some implementation details

- **How does local name server know root servers?**

- Need to configure name server with *root cache* file
- Contains root name servers and their addresses

```
.                3600000  NS   A.ROOT-SERVERS.NET.
A.ROOT-SERVERS.NET. 3600000  A    198.41.0.4
.                3600000  NS   B.ROOT-SERVERS.NET.
B.ROOT-SERVERS.NET. 3600000  A    128.9.0.107
...
```

- **How do you get addresses of other name servers**

- To lookup names ending `.stanford.edu.`, ask `Argus.stanford.edu.`
- Chicken and egg problem:
How to get `Argus.stanford.edu.`'s address?
- Solution: **glue** records – A records in parent zone
- Name servers for `edu.` have A record of `Argus.stanford.edu.`

Glue Record Example

- **Look up** `www.scs.stanford.edu` **assuming no cache**

```
dig +norec www.scs.stanford.edu @a.root-servers.net
```

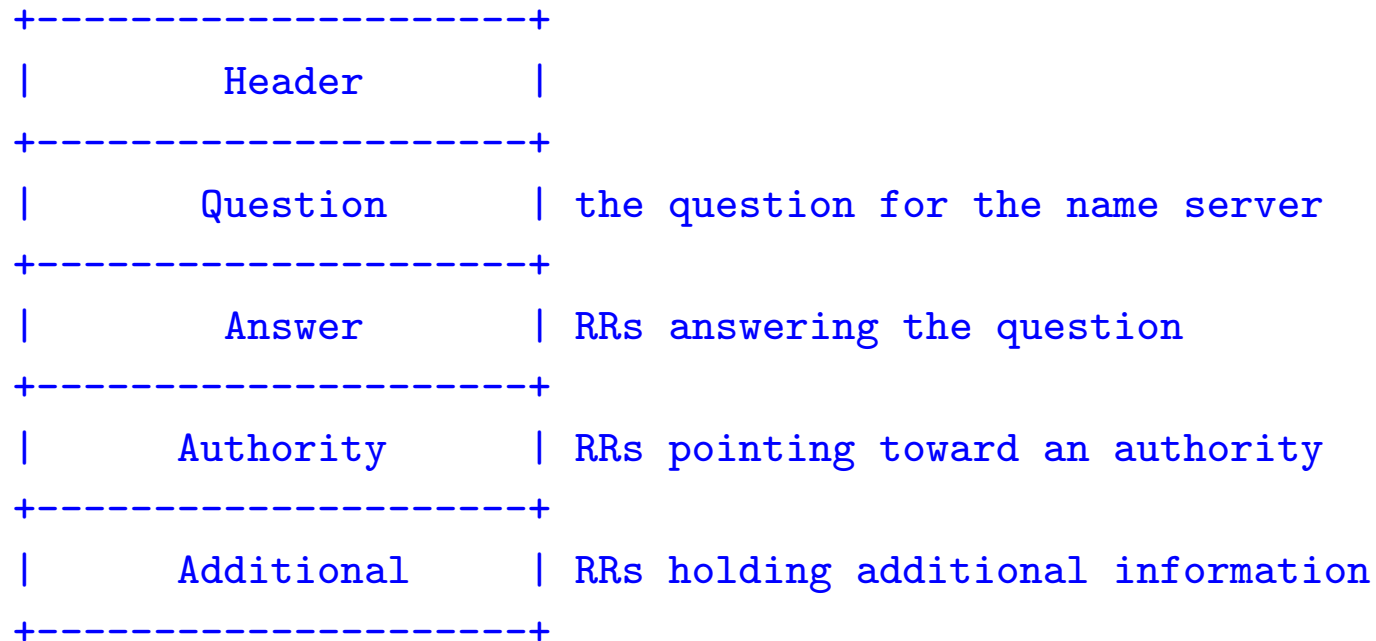
```
dig +norec www.scs.stanford.edu @a.gtld-servers.net
```

```
dig +norec www.scs.stanford.edu @argus.stanford.edu
```

```
dig +norec www.scs.stanford.edu @mission.scs.stanford.edu
```

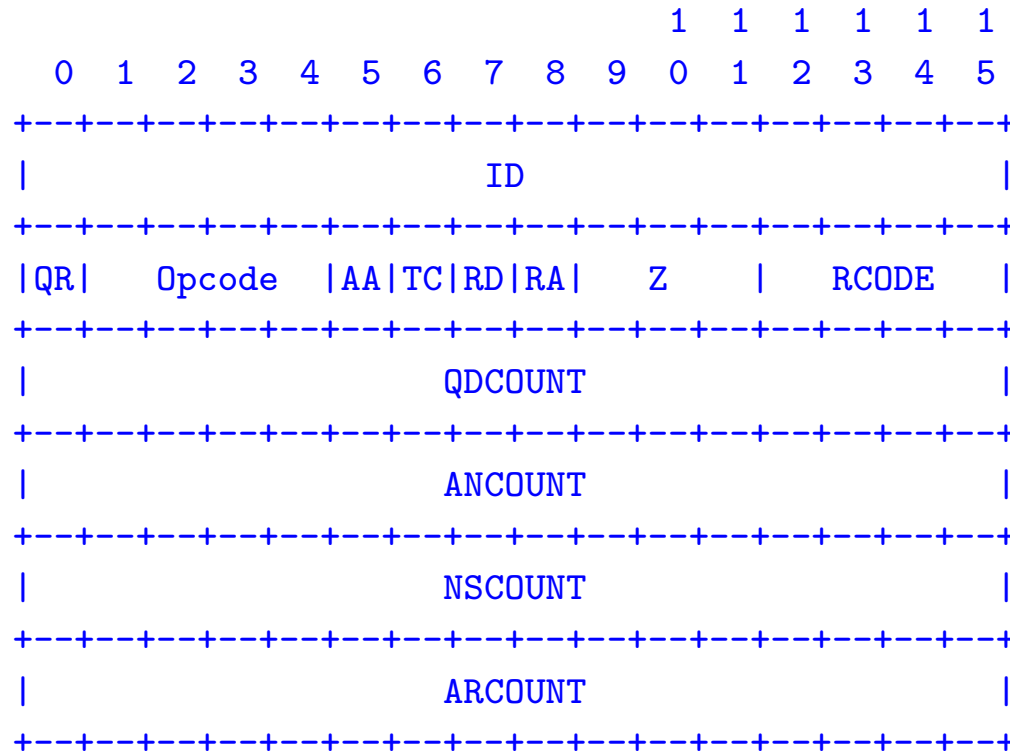
- **Get intermediary results for** `.edu`, `stanford.edu`, `scs.stanford.edu`, **and** `www.scs.stanford.edu`
- **Where are the glue records?**

Structure of a DNS message [RFC 1035]



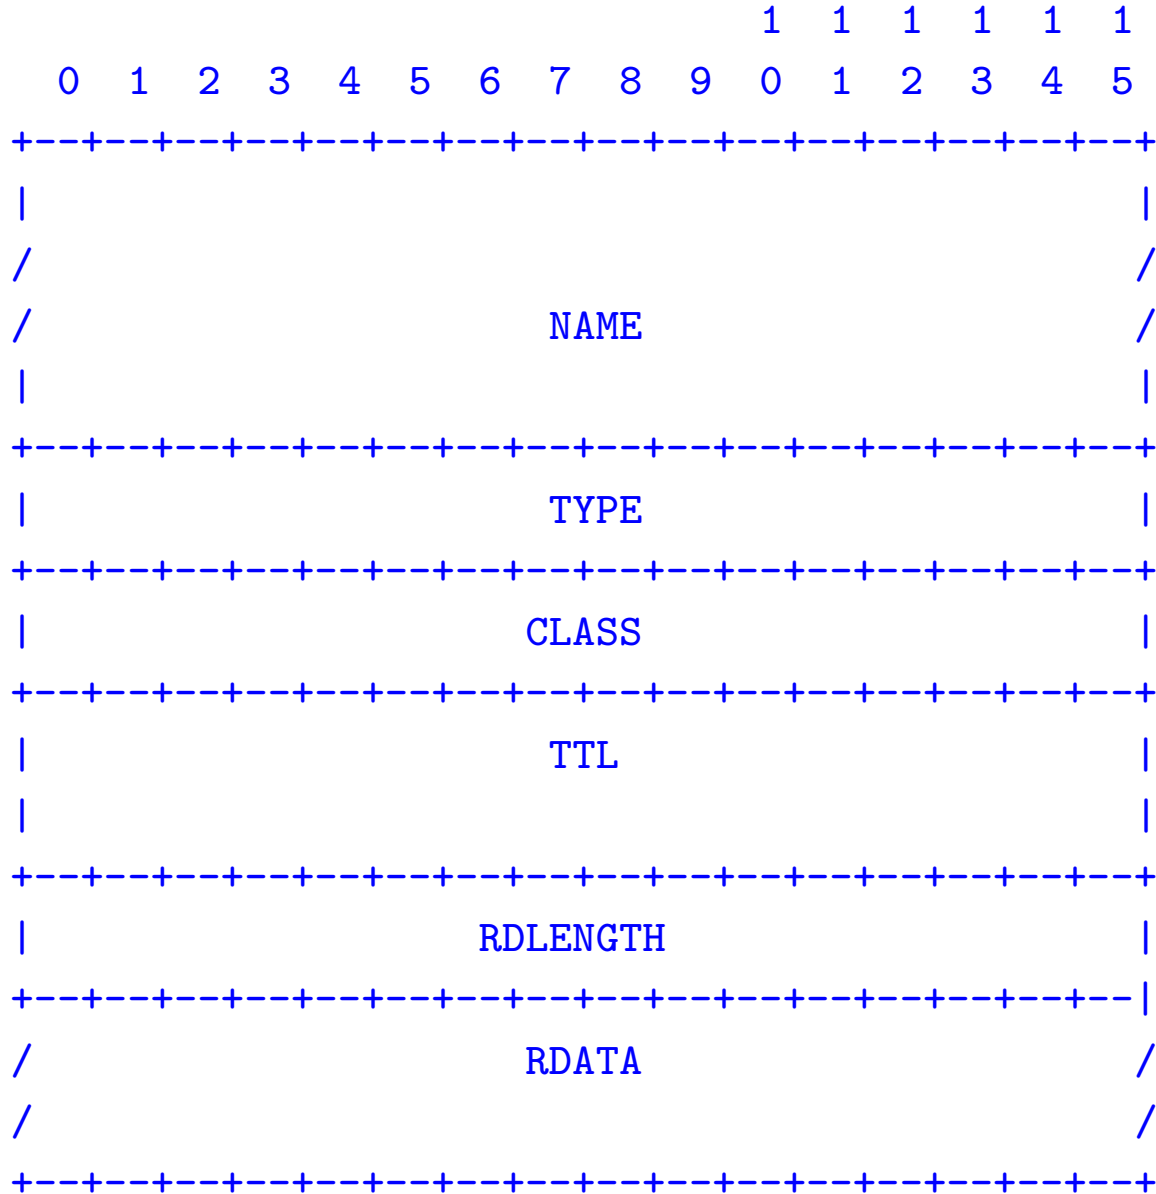
- **Same message format for queries and replies**
 - Query has zero RRs in Answer/Authority/Additional sections
 - Reply includes question, plus has RRs
- **Authority allows for delegation**
- **Additional for glue + other RRs client might need**

Header format



- **QR** – 0=query, 1=response
- **RCODE** – error code
- **AA**=authoritative answer, **TC**=truncated,
RD=recursion desired, **RA**=recursion available


Encoding of RRs



Encoding of domain names

- **A DNS name consists of a series of labels**
 - `www.stanford.edu.` has 3 labels: `www`, `stanford`, and `edu`
 - Labels can contain letters, digits, and “-”, but should not start or end with “-”
 - Maximum length 63 characters
 - Encoded as length byte followed by label
 - Last label always empty label
- **Names are case insensitive**
 - But server must preserve case of question in replies
 - Example: request `www.sTANford.EDu`, look at authority

Name compression



- **Observation: many common suffixes in DNS messages**
 - Particularly because of case preservation rule
- **Allow pointer labels to re-use suffixes**
 - Recal label starts with length byte (0-63)
 - If value $\geq 0xc0$ (192), subtract $0xc000$ from first *two* bytes, and treat as pointer into message

Secondary servers

- **Availability requires geographically disparate replicas**
 - E.g., I ask MIT to serve `scs.stanford.edu`
- **Typical setup: One master many slave servers**
- **How often to sync up servers? Trade-off**
 - All the time \implies high overhead
 - Rarely \implies stale data
- **Put trade-off under domain owner's control**
 - Fields in SOA record control secondary's behavior
 - Primary can change SOA without asking human operator of secondary
 - Primary can also give secondary a hint to check freshness

Other Records

- **Start of Authority (SOA) record**
 - States administrative information for a zone
 - `dig stanford.edu soa`
 - `dig sing.stanford.edu ns`
 - Tells you how long you can cache negative results
- **Mail Exchange (MX) record**
 - For historical reasons, mail does not have to use A records directly
 - Example: `ping scs.stanford.edu`
 - No such host, but you can still mail CS144 staff there

CNAME records

- CNAME record specifies an alias:

name [TTL] [IN] CNAME *canonical-name*

- As if any RR's associated w. *canonical-name* also for *name*
- Can look up with AI_CANONNAME flag to getaddrinfo

- Examples, to save typing:

```
wb.scs.stanford.edu.  CNAME  williamsburg-bridge.scs.stanford.edu.  
mb.scs.stanford.edu.  CNAME  manhattan-bridge.scs.stanford.edu.
```

- CNAME precludes any other RRs for name

- E.g., might want: david.com CNAME david.stanford.edu
- Illegal, because david.com would need NS records

- Note answer section can have CNAME for query
name + other RR(s) for *canonical-name*

- But don't point MXes to CNAMEs, as no A recs in additional section (try bad-mx.scs.stanford.edu.)

Reverse Lookups

- Remember traceroute...
- Traceroute can learn names of hosts through *reverse lookup*
- 128.30.2.121 → 121.2.30.128.in-addr.arpa
- PTR record points to canonical name
- Example: tinyos.stanford.edu → sing.stanford.edu
→ 65.76.67.171.in-addr.arpa ptr

Mapping addresses to names

- PTR records specify names

name [TTL] [IN] PTR "*ptrdname*"

- *name* – somehow encode address...how?
- *ptrdname* – domain name for this address

- IPv4 addrs stored under in-addr.arpa domain

- Reverse name, append in-addr.arpa
- To look up 171.66.3.9 → 9.3.66.171.in-addr.arpa.
- Why reversed? Delegation!

- IPv6 under ip6.arpa

- Historical note: ARPA funded original Internet

Using DNS for load-balancing

- **Can have multiple RR of most types for one name**
 - Required for NS records (for availability)
 - Useful for A records
 - (Not legal for CNAME records)
- **Servers rotate order in which records returned**
 - getaddrinfo returns a linked list of addrinfo structures
 - Most apps just use first address returned
 - Even if your name server caches results, clients will be spread amongst servers
- **Example: dig cnn.com multiple times**

SRV records

- Service location records

_service._proto.name [...] SRV prio weight port target

- **_service** – E.g., `_sfs` for NYU's SFS file system
- **_proto** – `_tcp` or `_udp`
- **name** – domain name record applies to
- **prio** – as with MX records, lower # → higher priority
- **weight** – within priority, affects randomization of order
- **port** – TCP or UDP port number (particularly useful for SIP)
- **target** – Server name, for which client needs A record

- Like a generalization of MX records for arbitrary services

DNS redirection for content distribution

- **Play with akamai and `www.microsoft.com`**

Classless in-addr delegation

- How to delegate on non-byte boundary?
- Solution: Use CNAME records
 - So-called *classless* in-addr delegation
- Example:

1.3.66.171.in-addr.arpa. CNAME 1.ptr.your-domain.com.

2.3.66.171.in-addr.arpa. CNAME 2.ptr.your-domain.com.

3.3.66.171.in-addr.arpa. CNAME 3.ptr.your-domain.com.

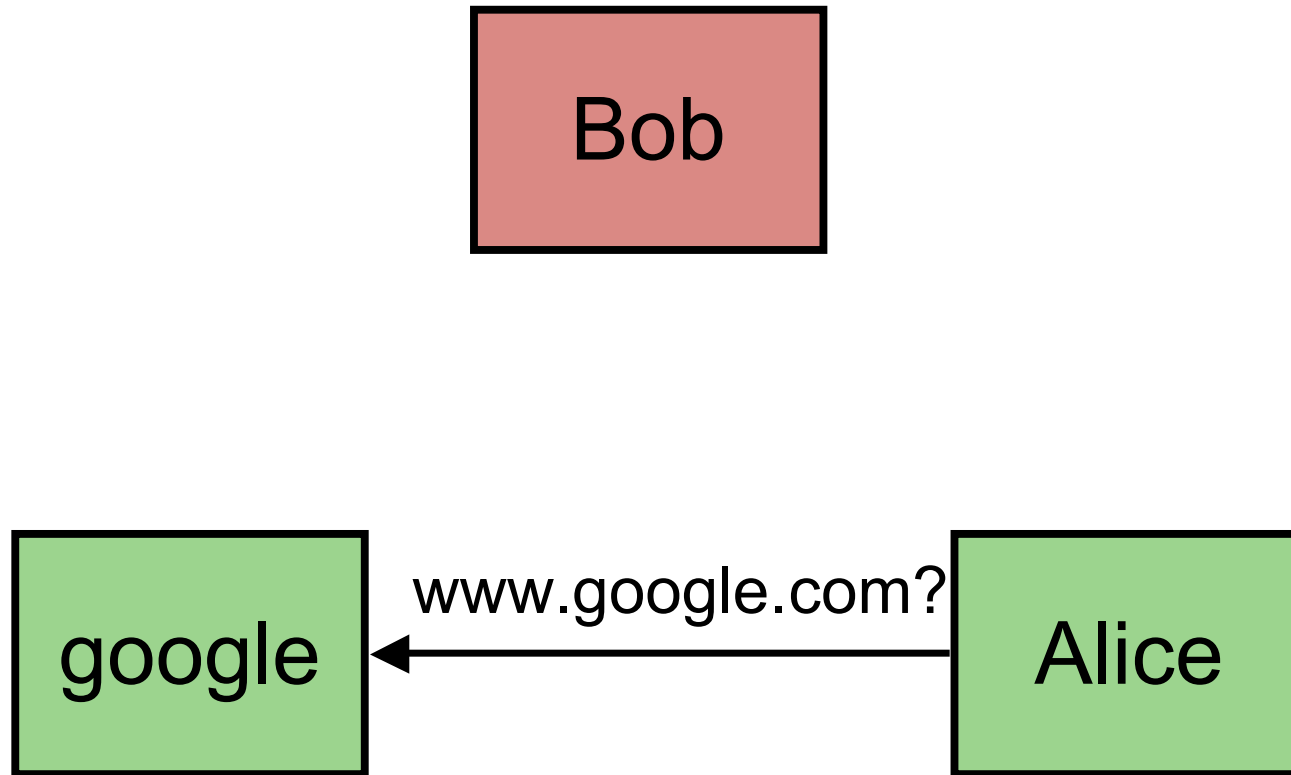
DNS exploits

- July 29, 2008, **Bruce Schneier**: “Despite the best efforts of the security community, the details of a critical internet vulnerability discovered by Dan Kaminsky about six months ago have leaked.”
- One of the basic problems: DNS caching
 - If you can poison the cache, the damage stays
 - Who knows how far it spreads...


DNS exploit example

- Alice wants to look up `www.google.com`
- Bob the attacker knows
- Bob knows source address/port, destination address/port
- Bob generates a spoof response: `www.google.com` is `www.evil.com`
- Challenge: Bob has to guess Query ID
- If Bob guesses, RR can stay in Alice's cache a long time

Exploit Example



Exploit Example




Countermeasures

- Choose good QIDs (used to be incremented, now randomly generated), 16 bits
- Randomize source port, 16 bits
- Some protection, but only makes it take longer, networks are faster each day

Another exploit

- **DNS clients used to trust all responses**
- **Problem: glue records and helpful A records**
 - Ask NS of `evil.com` for `www.evil.com`
 - Says `www.evil.com` is a CNAME for `www.google.com`
 - Provides A record for `www.google.com`

Exploit Example



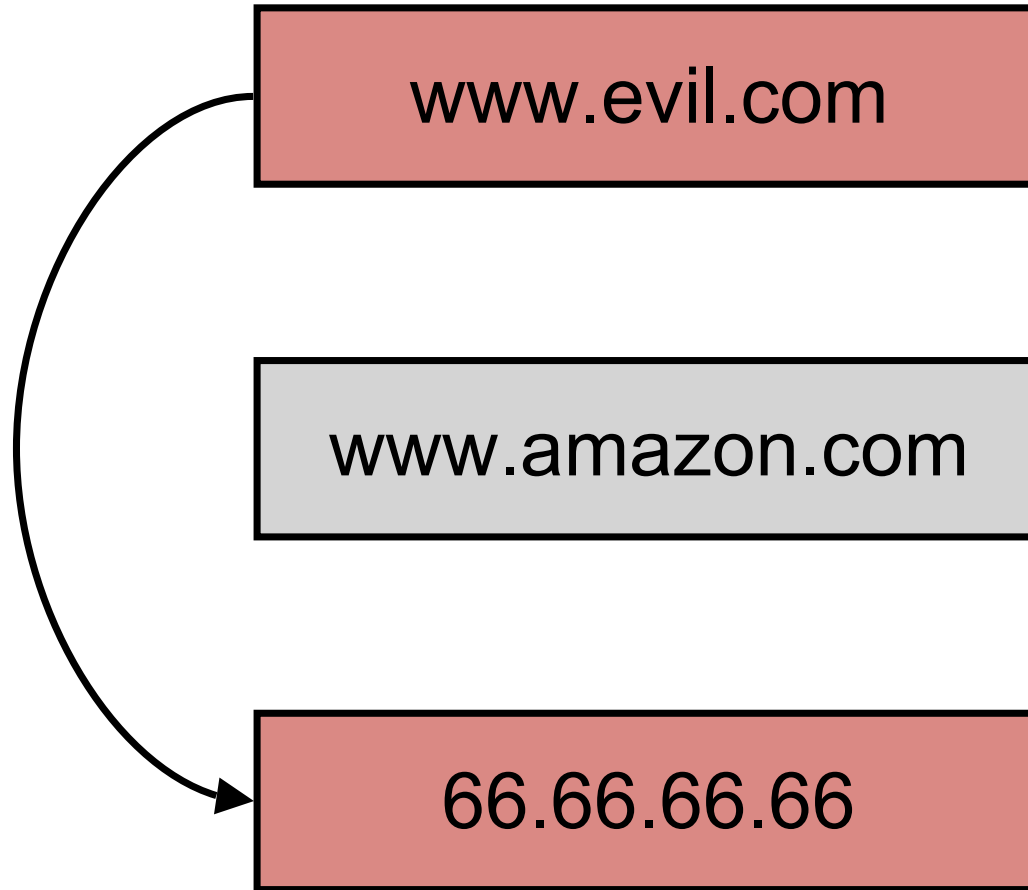
It gets worse

- **Glue records can overwrite standard A records**
- **Even if you have a good A record for `www.amazon.com`, it's overwritten**
- **E.g., Server wants name of my IP address**
 - Looks up `66.66.66.66.in-addr.arpa`
- **I say nameserver for `66.66.66.66.in-addr.arpa` is `www.amazon.com`**
 - Include glue A record for `www.amazon.com` in my reply

Solution 1

- **Only use glue records for duration of query**
 - Cache only end-to-end traversal of pointers, not intermediate steps
- **In CNAME example `www.evil.com` will point to evil server**
 - `www.amazon.com` will not point to evil server
- **In `in-addr.arpa` example, can lie about hostname**
 - But I can lie anyway
 - Have to check reverse lookup result by doing forward lookup

Example



Solution 2: bailiwick checking

- **Only pay attention to answers for the domain you've asked**
- **Response from `evil.com` can't tell you the A record for `google.com`**
- **Ask `google.com` for `www.google.com`**
- **Opponent can still race, but at least it's not deterministic**

Recent Kaminsky exploit

- Make winning the race easier
- Brute force attack
- Force Alice to look up AAAA.google.com, AAAB.google.com, etc.
- Forge CNAME responses for each lookup, inserting A record for www.google.com
- Circumvents bailiwick checking

Solution: signatures

- **Signature: cryptographic way to prove a party is who they say they are (more later in quarter)**
- **Requires a chain of trust**
- **Whom do you trust to sign DNS?**

DNS Overview

- Distributed system for mapping names to values (e.g., IP addresses)
- Read-dominated workload allows caching
- Name structure allows distribution, independent administration
- Caching means bad data can stay a long time
- Standard protocol does not authenticate response is from server: DNSSec does