

# **EECS 498 – Lecture Notes #1b**

## **Introduction to Distributed Systems**

**Farnam Jahanian**  
**Department of EECS**  
**University of Michigan**

**EECS 498 Lecture Notes**

**<http://www.eecs.umich.edu/~farnam>**

### **Lectures: Weeks 1-3**

- Introduction to Networking: communication issues and network technologies  
*Reading list:* Tanenbaum text Chapter 2.1, pp 57-68, Layered Protocols  
Handout in class: Chapter 5: -- “Protocols Underlying HTTP” from “Web Protocols and Practices” by Krishnamurthy & Rexford, 2001.
- Introduction to distributed systems, characteristics of distributed systems, design issues, h/s concepts, distributed programming models  
*Reading list:* Tanenbaum text Chapter 1, pp 1-42.
- Overview of Distributed Computing Paradigms  
*Reading List:*  
Handout in class: Chapter 3 – from “Distributed Computing: Principles and Applications” text by M.L. Liu (recommended text)

# Introduction to Distributed Systems

## Distributed Systems

---

### Three Technology Advances:

- Development of powerful microprocessors
- Development of high-speed networks
- Development of denser and cheaper memory/storage

**Easy:** put together large # of powerful processors connected by a high-speed network.

**Hard:** SOFTWARE! SOFTWARE! SOFTWARE!

# Distributed Systems

---

## What is a distributed system?

*"You know you have one when the crash of a computer you've never heard of stops you from getting any work done."* Leslie Lamport

- A collection of (perhaps) heterogeneous nodes connected by one or more interconnection networks which provides access to system-wide shared resources and services.
- A collection of independent computers that appears to its users as a single coherent system.


## Examples?

# Characteristics of a distributed systems

---

- **Multiple Computers:**  
More than one physical computer, each consisting of CPUs, local memory, and possibly stable storage, and I/O paths to connect it with the environment.
- **Interconnections:**  
Mechanisms for communicating with other nodes via a network.
- **Shared State:**  
If a subset of nodes cooperate to provide a service, a shared state is maintained by these nodes. The shared state is distributed or replicated among the participants.

## An Abstract View Distributed Systems



A distributed system organized as middleware.  
Note that the middleware layer extends over multiple machines.

## Distributed vs. Centralized Systems

Why distribute?

- Resource sharing
- Device sharing
- Flexibility to spread load
- Incremental growth
- Cost/performance
- Reliability/Availability
- Inherent distribution
- Security?

---

### **Why NOT distribute?**

- Software
- Network
- Security
- System management

### **Numerous sources of complexity including:**

- Transparent/uniform access to data or resources
- Independent failures of processors (or processes)
- Dynamic membership of a system
- Unreliable/unsecured communication
- Overhead in message communication
- Distribution or replication of data (or meta-data)
- Lack of clean common interfaces

## **Design Goals & Issues**

---

- Connecting users and resources is the primary goal
- Transparency: hide the fact that processes and resources are physically distributed
- Openness: offer services according to rules and interfaces that describe the syntax and semantics of those services
  - Interoperability and portability
  - Separating policy from mechanism
- Scalability
- Performance
- Dependability

## Transparency in a Distributed System

Transparency	Description
Access	Hide differences in data representation and how a resource is accessed
Location	Hide where a resource is located
Migration	Hide that a resource may move to another location
Relocation	Hide that a resource may be moved to another location while in use
Replication	Hide that a resource may be shared by several competitive users
Concurrency	Hide that a resource may be shared by several competitive users
Failure	Hide the failure and recovery of a resource
Persistence	Hide whether a (software) resource is in memory or on disk

Different forms of transparency in a distributed system.

## Transparency

- How to achieve “single-system image”? How to hide distribution from users or programs?
- Is it a good idea?

Sometimes requires trade off transparency for performance

## Scalability

- The challenge is to build distributed systems that scale with the increase in the number of CPUs, users, and processes, larger databases, etc.
- Scalability along several dimensions: size, geography, administrative domains

## Scalability Problems

Concept	Example
<b>Centralized services</b>	<b>A single server for all users</b>
<b>Centralized data (tables)</b>	<b>A single on-line telephone book</b>
<b>Centralized algorithms</b>	<b>Doing routing based on complete information</b>

Examples of scalability limitations.

## How to scale?

---

A very simple principle:

Avoid centralized services, centralized tables, and centralized algorithms

Characteristics of decentralized algorithms:

- No machine has complete information about the system state
- Machines make decisions based only on local information
- Failure of one machine does not ruin the algorithm
- There is no implicit assumption about a global clock

---

A few lessons from AFS:

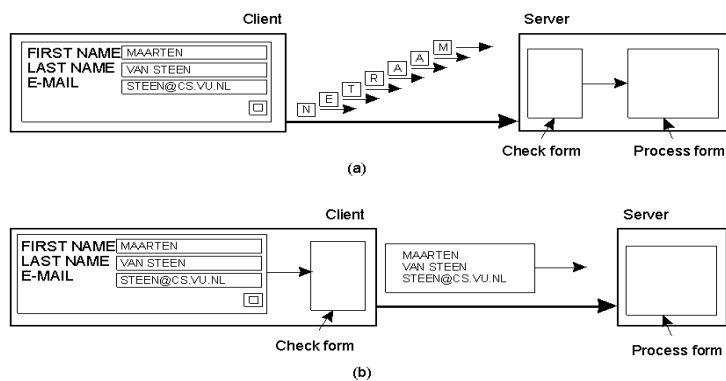
- “Clients have cycles to burn.”
- “Cache whenever possible.”
- “Exploit usage properties.”
- “Minimize system-wide knowledge/change.”
- “Batch if possible.”
- Multicast often works!



## Scaling Techniques (Tanenbaum's Text)

- Hiding communication latencies
  - ❖ Asynchronous communication
  - ❖ Function shipping to clients
- Distribution of components
  - ❖ DNS name space
- Caching and Replication
  - ❖ Maintaining consistency

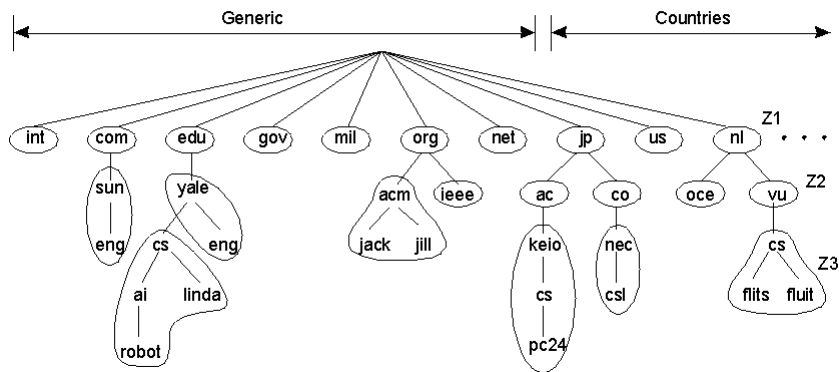
## Scaling Techniques (1)



a) a server or

b) a client check forms as they are being filled

## Scaling Techniques (2)



An example of dividing the DNS name space into zones.

## Performance

Various performance metrics:

- response time
- throughput
- system utilization
- network capacity utilization

Key issue in parallelizing computations in a distributed system?

*overhead of message communication*

## Performance

---

Trade off:

- More tasks  $\rightarrow$  more parallelism  $\rightarrow$  better performance
- More tasks  $\rightarrow$  more communication  $\rightarrow$  worse performance

Grain size affects # of messages:

fine-grained parallelism	vs.	coarse-grained parallelism
small computations		large computations
high interaction rate		low interaction rate

## Dependability

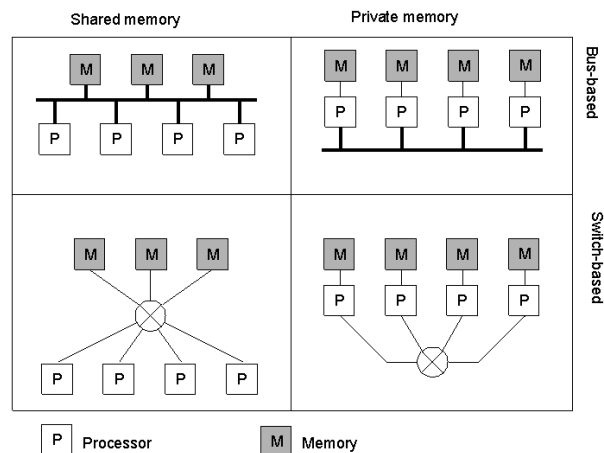
---

- **Reliability:** measured by the probability  $R(t)$  that the system is up (and providing correct service) during the time interval  $[0, t]$  assuming that it was operational at time  $t$ .
- **Availability:** measured by the probability  $A(t)$  that the system is operational at the instant of time  $t$ . As  $t \rightarrow \infty$ , availability expresses the fraction of time that the system is usable.
- **Timeliness:** ability to meet timing constraints imposed on task execution or service delivery.
- **Integrity:** replicated copies of data must be kept *consistent*.
- **Security:** protection from unauthorized usage/access. Why more difficult in distributed systems?

## Distributed Programming Paradigms

- Client/server model
- Remote procedure calls
- Distributed File Systems
- Group communication and multicasts
- Distributed transactions
- Distributed shared memory
- Distributed object-based systems
- Publish-subscribe model
- Peer-to-peer model
- The Web


## Hardware Concepts



Different basic organizations and memories in distributed computer systems


## Multiprocessors (1)

- A bus-based multiprocessor.




## Multiprocessors (2)

- A crossbar switch
- An omega switching network



(a)




(b)


## Homogeneous Multicomputer Systems

---

- a) Grid
- b) Hypercube



(a)



(b)

## Heterogeneous Multicomputer Systems

---

- Most distributed systems today are built on top of heterogeneous multicomputers and interconnection networks
- No global system view
- Sophisticated software needed to support distributed applications

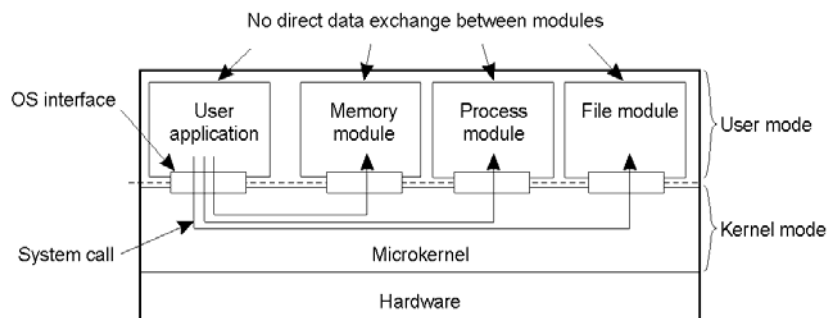
## Software Concepts

System	Description	Main Goal
<b>DOS</b>	Tightly-coupled operating system for multi-processors and homogeneous multicomputers	Hide and manage hardware resources
<b>NOS</b>	Loosely-coupled operating system for heterogeneous multicomputers (LAN and WAN)	Offer local services to remote clients
<b>Middleware</b>	Additional layer atop of NOS implementing general-purpose services	Provide distribution transparency

- An overview of
- **DOS (Distributed Operating Systems)**
- **NOS (Network Operating Systems)**
- **Middleware**

## Uniprocessor Operating Systems

- Separating applications from operating system code through a microkernel.



## Multiprocessor Operating Systems (1)

```
monitor Counter {  
    private:  
        int count = 0;  
    public:  
        int value() { return count;}  
        void incr () { count = count + 1;}  
        void decr() { count = count - 1;}  
}
```

- A monitor to protect an integer against concurrent access.

## Multiprocessor Operating Systems (2)

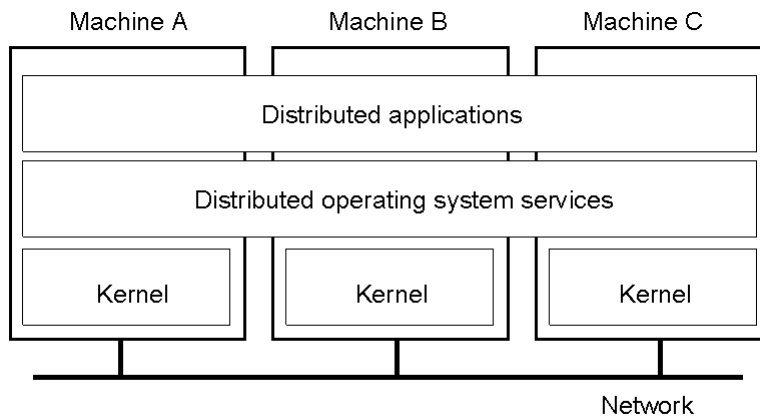
```
monitor Counter {  
    private:  
        int count = 0;  
        int blocked_procs = 0;  
        condition unblocked;  
    public:  
        int value () { return count;}  
        void incr () {  
            if (blocked_procs == 0)  
                count = count + 1;  
            else  
                signal (unblocked);  
        }  
        void decr() {  
            if (count ==0) {  
                blocked_procs = blocked_procs + 1;  
                wait (unblocked);  
                blocked_procs = blocked_procs - 1;  
            }  
            else  
                count = count - 1;  
        }  
}
```

- A monitor to protect an integer against concurrent access, but blocking a process.




## Multicomputer Operating Systems (1)

- General structure of a multicomputer operating system



## Multicomputer Operating Systems (2)

- Alternatives for blocking and buffering in message passing.



## Multicomputer Operating Systems (3)

Synchronization point	Send buffer	Reliable comm. guaranteed?
Block sender until buffer not full	Yes	Not necessary
Block sender until message sent	No	Not necessary
Block sender until message received	No	Necessary
Block sender until message delivered	No	Necessary

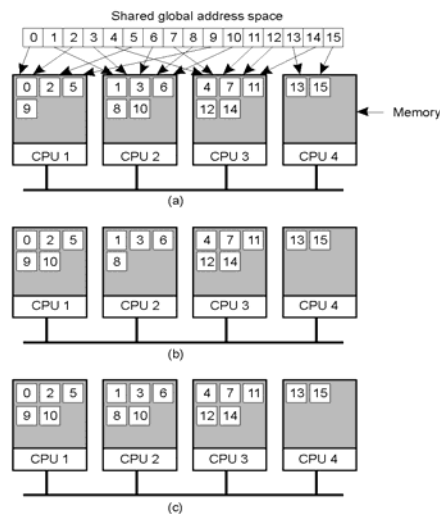
- Relation between blocking, buffering, and reliable communications.

## Distributed Shared Memory Systems (1)

- a) **Pages of address space distributed among four machines**

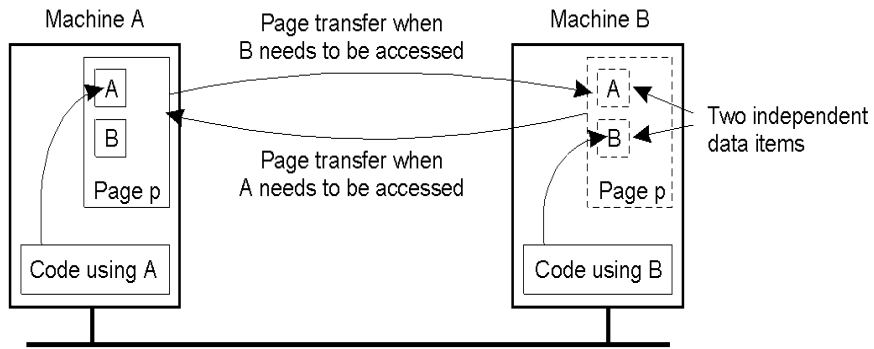
- b) **Situation after CPU 1 references page 10**

- c) **Situation if page 10 is read only and replication is used**




## Distributed Shared Memory Systems (2)

- False sharing of a page between two independent processes.



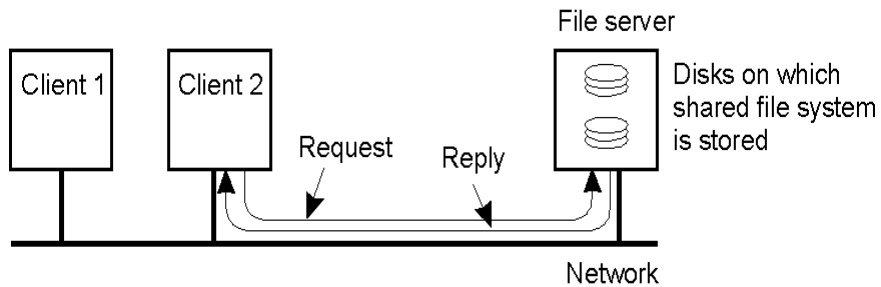
## Network Operating System (1)

- General structure of a network operating system.



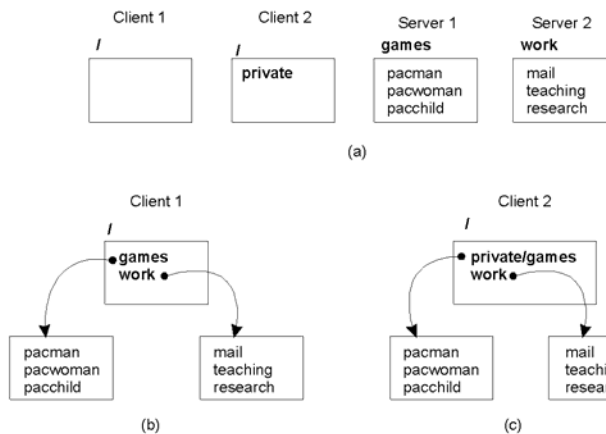
## Network Operating System (2)

- Two clients and a server in a network operating system.




## Network Operating System (3)

- Different clients may mount the servers in different places.




## Positioning Middleware

- General structure of a distributed system as middleware.



## Middleware and Openness



- In an open middleware-based distributed system, the protocols used by each middleware layer should be the same, as well as the interfaces they offer to applications.

## Comparison between Systems

- A comparison between multiprocessor operating systems, multicomputer operating systems, network operating systems, and middleware based distributed systems.

Item	Distributed OS		Network OS	Middleware-based OS
	Multiproc.	Multicomp.		
<b>Degree of transparency</b>	<b>Very High</b>	<b>High</b>	<b>Low</b>	<b>High</b>
<b>Same OS on all nodes</b>	<b>Yes</b>	<b>Yes</b>	<b>No</b>	<b>No</b>
<b>Number of copies of OS</b>	<b>1</b>	<b>N</b>	<b>N</b>	<b>N</b>
<b>Basis for communication</b>	<b>Shared memory</b>	<b>Messages</b>	<b>Files</b>	<b>Model specific</b>
<b>Resource management</b>	<b>Global, central</b>	<b>Global, distributed</b>	<b>Per node</b>	<b>Per node</b>
<b>Scalability</b>	<b>No</b>	<b>Moderately</b>	<b>Yes</b>	<b>Varies</b>
<b>Openness</b>	<b>Closed</b>	<b>Closed</b>	<b>Open</b>	<b>Open</b>