

(1) FEATURE EXPECTATIONS [5 min]

1. Use cases
2. Scenarios that will **not** be covered
3. Who will use
4. How many will use
5. Usage patterns

(2) ESTIMATIONS [5 min]

1. Throughput (QPS **for Read and Write** queries)
2. Latency expected **from** the system (**for read and write** queries)
3. **Read/Write** ratio
4. Traffic estimates
 - **Write** (QPS, Volume **of** data)
 - **Read** (QPS, Volume **of** data)
5. Storage estimates
6. Memory estimates
 - **If** we are **using** a cache, what **is** the kind **of** data we want **to** store **in** the cache
 - How much RAM **and** how many machines **do** we need **for** us **to** achieve this?
 - Amount **of** data you want **to** store **in** disk/SSD

(3) DESIGN GOALS [5 min]

1. Latency and Throughput requirements
2. Consistency vs Availability [Weak/strong/**eventual** => consistency | Failover/**replication** => availability]

(4) HIGH-LEVEL DESIGN [5-10 min]

1. APIs **for Read/Write** scenarios **for** crucial components
2. **Database schema**
3. Basic algorithm
4. High-level design **for Read/Write** scenario

(5) DEEP DIVE [15-20 min]

1. Scaling the algorithm
2. Scaling individual components:
 - Availability, Consistency and Scale story for each component
 - Consistency and availability patterns
3. Think about the following components, how they would fit in and how it would help
 - a) DNS
 - b) CDN [Push vs Pull]
 - c) Load Balancers [Active-Passive, Active-Active, Layer 4, Layer 7]
 - d) Reverse Proxy
 - e) Application layer scaling [Microservices, Service Discovery]
 - f) DB [RDBMS, NoSQL]
 - RDBMS
 - Master-slave, Master-master, Federation, Sharding, Denormalization, SQL Tuning
 - NoSQL
 - Key-Value, Wide-Column, Graph, Document
 - Fast-lookups:
 - RAM [Bounded size] => Redis, Memcached
 - AP [Unbounded size] => Cassandra, RIAK, Voldemort
 - CP [Unbounded size] => HBase, MongoDB, Couchbase, DynamoDB
 - g) Caches
 - Client caching, CDN caching, Web server caching, Database caching, Application caching, Cache @Query level, Cache @Object level
 - Eviction policies:
 - Least Recently Used (LRU)
 - Least Frequently Used (LFU)
 - First in First Out (FIFO)
 - Cache Loading Policies
 - Cache aside
 - Write through
 - Write behind
 - Refresh ahead
 - h) Asynchronism
 - Message queues
 - Task queues
 - Backpressure
 - i) Communication
 - TCP, UDP, REST, RPC, Thrift, GraphQL

(6) JUSTIFY [5 min]

1. Throughput of each layer
2. Latency caused between each layer
3. Overall latency justification